

# FaaSFlow: Enable Efficient Workflow Execution for Function-as-a-Service

(W6) Paper Reading

Nan Lin

Shanghai Jiao Tong University

2024-08-11

# FaaSFlow: Enable Efficient Workflow Execution for Function-as-a-Service

Zijun Li

Shanghai Jiao Tong University  
Shanghai, China  
lzjzx1122@sjtu.edu.cn

Yushi Liu

Shanghai Jiao Tong University  
Shanghai, China  
ziliuziliulys@sjtu.edu.cn

Linsong Guo

Shanghai Jiao Tong University  
Shanghai, China  
gls1196@sjtu.edu.cn

Quan Chen

Shanghai Jiao Tong University  
Shanghai, China  
chen-quan@cs.sjtu.edu.cn

Jiagan Cheng

Shanghai Jiao Tong University  
Shanghai, China  
chengjiagan@sjtu.edu.cn

Wenli Zheng

Shanghai Jiao Tong University  
Shanghai, China  
zheng-wl@cs.sjtu.edu.cn

Minyi Guo

Shanghai Jiao Tong University  
Shanghai, China  
guo-my@cs.sjtu.edu.cn

- ASPLOS '22, February 28 - March 4, 2022
- Keywords: *Decentralization*

# Outline

**Review**

**Motivation**

**Architecture**

**Evaluation**

# Outline

## Review

## Motivation

## Architecture

## Evaluation

# Serverless Workflow

## Serverless Workflow

Serverless functions are event-driven, and they need to be executed in a pre-defined order. Such a diagram with nodes connected by edges in a DAG form is known as the **serverless workflow**.




Figure 2: The example DAG-based workflows.

# Serverless Workflow

## Control-Plane and Data-Plane

- **Control-Plane**: User-defined execution order
- **Data-Plane**: Runtime data dependency

Usually identical and static. However, in serverless context, auto-scaling and warm containers may lead to multiple and different scales in the data-plane.

# Outline

Review

**Motivation**

Architecture

Evaluation

# MasterSP and its Limitations

## Master-side Workflow Schedule Pattern

**Centralized** Workflow: Central workflow engine in the master node determines whether a function task is triggered to run or not.

- Engine makes resource provision
- Task  $T_f$  triggered only if its predecessors are all completed:
  1. Assign  $T_f$  from the master engine
  2. Execute  $T_f$  invocation
  3. Return the execution state to the master engine




Figure 1: The overhead analysis for traditional workflow execution architecture in serverless context. **WorkerSP**

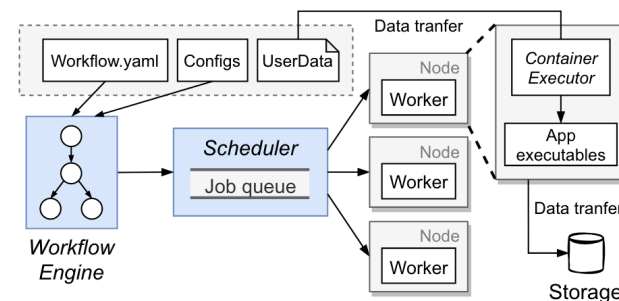



Figure 3: Implemented prototype of HyperFlow-serverless.



# MasterSP and its Limitations

Problems:

- Large scheduling overhead: Transfer of function execution states
- Large data movement overhead: Additional database storage services for temp data storage and delivery



**Figure 4: The scheduling overhead of executing the workflow benchmark (the scheduling overhead and the end-to-end latency depend on the critical path).**

# Data-Shipping Pattern

## Data-Shipping Pattern

Each time a function task runs, the input data needs to be fetched from its predecessor functions, then read into memory for execution by the container executor. Such process is called a **data-shipping pattern**.

Problems:

- Function isolation brings more overhead of task-to-task data communication
- Compulsory for user to use remote storage services
- Data locality is not utilized

# Outline

Review

Motivation

**Architecture**

Evaluation

# WorkerSP's Structure Organization

The inverse of MasterSP: **Worker-side workflow schedule pattern (WorkerSP)**, *Decentralize*

## Structure Organization of WorkerSP

- Master node scheduling  $\xrightarrow{\text{Offload}}$  Per-worker engine assigned to perform local function triggering and invoking
- Master node only partition a workflow graph into sub-graph (See later)
- *Workflow* structure introduced with *State*, *FunctionInfo* and *InvocationID*.
  - *State*: Execution state of functions and their predecessors for invocation synchronization
  - *FunctionInfo*: Meta information for local functions
  - *InvocationID*: Unique state identification

# WorkerSP's State Synchronization

(Reminder) Engine of each worker node maintains functions' and their predecessors' execution state in the **local sub-graph**.

## Example: Invocation Synchronization

1.  $F_A$  is invoked
2. State pass to Node  $B$  and  $C$
3.  $F_B$  and  $F_C$  update info
4. When  $F_A$  finished, PredecessorsDone of  $B$  and  $C + 1$
5. When PredecessorsDone = PredecessorsCount, local engine of  $B$  and  $C$  will trigger

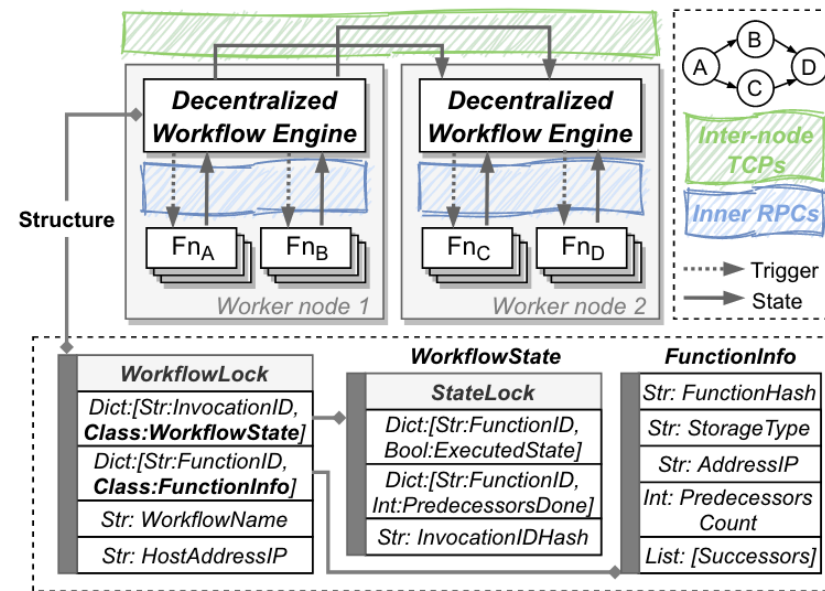


Figure 6: The data structure for workflow triggering and invocation management in the WorkerSP.

# Overview of FaaSFlow

## FaaSFlow: Workflow System

Three components:

1. **Workflow graph scheduler**
2. **Per-worker workflow engine**
3. Adaptive Storage Library **FaaStore**

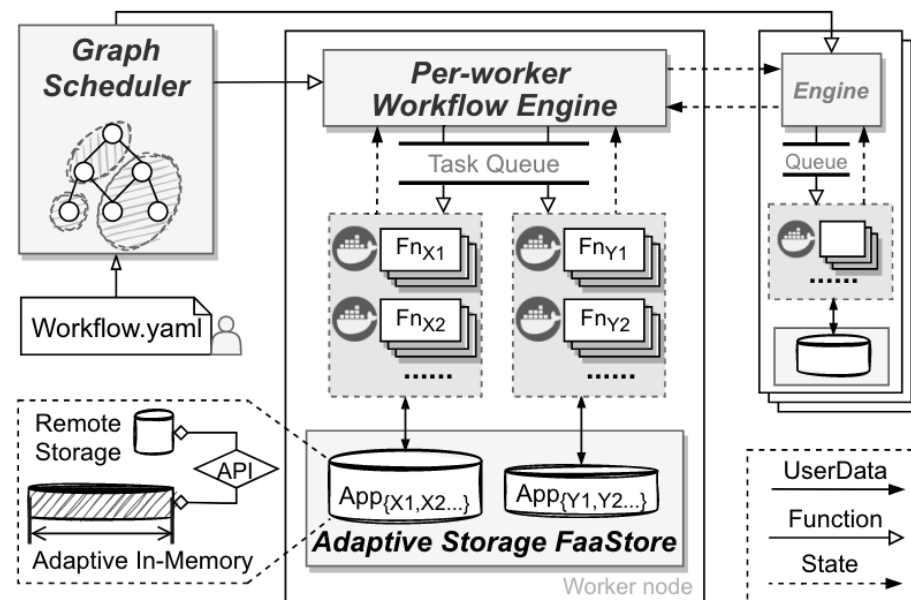



Figure 8: System design of FaaSFlow.

# Component 1: Graph Scheduler

1. **DAG Parser** parse the hierarchy *Workflow Definition Language (WDL)* (which defines a serverless workflow)



**Figure 9: The supported logic flows in FaaSFlow and virtual nodes introduced in the parsing steps.**

# Component 1: Graph Scheduler

## 2. Graph Partitionning: Partitioning of DAG

To alleviate the gap between Control-plane and (dynamic data-plane):

- $\text{Scale}(v_i)$ : Avg. number of scaled instances of a function node  $v_i$  during iteration
- $\text{Map}(v_i)$ : Mapped instances in the data-plane (e.g. Foreach)

Partition iteration activated when significant performance degradation

**Algorithm 1:** Functions grouping and scheduling.

**Data:**  $\text{Cap}[\text{node}]$ : a list of the capacity of containers left to be created on each node.  $\text{Quota}(G)$ : the in-memory storage quota of graph  $G$ , discussed in Section 4.3.1.

**Input:** workflow graph  $G$ , functions  $f_1, f_2, \dots, f_n$

```

1  $S \leftarrow \{f_1\}, \{f_2\}, \dots, \{f_n\}$ ;  $W \leftarrow \text{RandomNodes}(S)$ ;
2  $\{f\}.\text{StorageType} \leftarrow \text{'DB'}$ ;  $\text{mem\_consume} \leftarrow 0$ ;
3 repeat
4    $\text{cpath} = (\{f\}, \{e\}) \leftarrow \text{critical\_path}(G)$ ;
5    $\text{descend\_sort\_by\_weight}(\{e\})$ ;  $\text{flag\_merge} \leftarrow \text{False}$ ;
6   for  $e$  in  $\{e\}$  do
7      $f_{\text{start}} \leftarrow e.\text{start}, f_{\text{end}} \leftarrow e.\text{end}$ ;
8      $S_{\text{start}} \leftarrow S[\text{S.find}(f_{\text{start}})]$ ,  $S_{\text{end}} \leftarrow S[\text{S.find}(f_{\text{end}})]$ ;
9     if  $S_{\text{start}} == S_{\text{end}}$  then continue; end
10     $n_{\text{start}} \leftarrow \sum_{s_i \in S_{\text{start}}} \text{Scale}(s_i)$ ,  $\text{Cap}[W(S_{\text{start}})] -= n_{\text{start}}$ ;
11     $n_{\text{end}} \leftarrow \sum_{s_i \in S_{\text{end}}} \text{Scale}(s_i)$ ,  $\text{Cap}[W(S_{\text{end}})] -= n_{\text{end}}$ ;
12    if  $n_{\text{start}} + n_{\text{end}} > \max(\text{Cap}[\text{node}])$  then continue; end
13    if  $f_{\text{start}}.\text{StorageType} == \text{'DB'}$  then
14      if  $\text{mem\_consume} + e.\text{weight} > \text{Quota}(G)$  then
15        continue; end
16       $\text{mem\_consume} += e.\text{weight}$ ;
17       $f_{\text{start}}.\text{StorageType} \leftarrow \text{'MEM'}$ ;
18    end
19    if  $(f_i, f_j) \subseteq S_{\text{start}} \cup S_{\text{end}}$  and  $(f_i, f_j) \notin \text{cont}(G)$  then
20       $S_{\text{new}} \leftarrow S_{\text{start}} \cup S_{\text{end}}$ ; else continue; end
21     $W[S_{\text{new}}] \leftarrow \text{binpack}(\text{limit} = \text{Cap}[\text{node}] > n_{\text{start}} + n_{\text{end}})$ ;
22     $\text{Cap}[W(S_{\text{new}})] += n_{\text{start}} + n_{\text{end}}$ ;
23     $S.\text{insert}(S_{\text{new}})$ ,  $S.\text{delete}(S_{\text{start}}, S_{\text{end}})$ ;
24     $\text{flag\_merge} \leftarrow \text{True}$ ; break;
25  end
26 until  $\text{flag\_merge} == \text{False}$ ;
27 return group  $S = \{S_1, S_2, \dots, S_k\}$ , Worker  $W[S_i] = \text{node}$ 

```

**Annotations:**

- 1: Each node as a group
- 2: Locate functions with longest edge
- 3: Merge
- 4: Do not exceed Max capacity
- 5: Do not exceed memory const.
- 6: Avoid memory contention functions




# Component 2: Per-Worker Workflow Engine

Maintaining states for different functions.

Direct state communication via **inter-node TCP** or **inner RPC connections**.

## Red-Black Deployment

Manage different versions of sub-graph versions in worker engines, only the up-to-date version is getting triggered.



(a) Red-Black Deployment

## Component 3: FaaStore

In-memory storage enables data and files reside in local sub-graph; defUlt remote store save them by user configs.

### In-Memory Quota

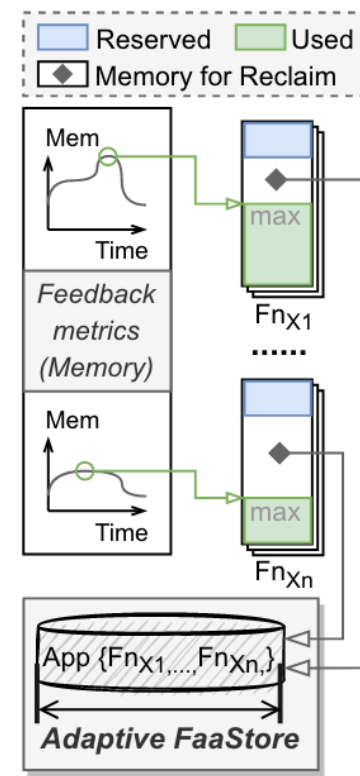
Well-organized quota for data movement

Due to *over-provisionning*,

$$O(v_i) = \max(\text{Mem}(v_i) - S - \mu, 0)$$

and

$$\text{Quota}(G(V, E)) = \sum_{v \in V} O(v)$$



(b) Memory Reclamation

# Outline

Review

Motivation

Architecture

**Evaluation**

# Evaluation

FaaSFlow reduces the scheduling overhead from 712ms to 141.9ms for scientific workflows, and from 181.3ms to 51.4ms for real-world applications on average. All applications can achieve an average of 74.6% scheduling overhead optimization in FaaSFlow.

*Basically did not mention memory allocation improvements .....*